

Mercedes-Benz Greener Manufacturing Project

In [1]:

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from sklearn.feature_selection import VarianceThreshold
from sklearn.preprocessing import LabelEncoder
```

In [4]:

```
# read in train and test data
data_train= pd.read_csv('mbtrain.csv')
data_test = pd.read_csv('mbtest.csv')

print(train.shape, test.shape)
```

(4209, 378) (4209, 377)

In [5]:

```
data_train.head()
```

Out[5]:

	ID	y	X0	X1	X2	X3	X4	X5	X6	X8	...	X375	X376	X377	X378	X379	X380
0	0	130.81	k	v	at	a	d	u	j	o	...	0	0	1	0	0	0
1	6	88.53	k	t	av	e	d	y	l	o	...	1	0	0	0	0	0
2	7	76.26	az	w	n	c	d	x	j	x	...	0	0	0	0	0	0
3	9	80.62	az	t	n	f	d	x	l	e	...	0	0	0	0	0	0
4	13	78.02	az	v	n	f	d	h	d	n	...	0	0	0	0	0	0

5 rows × 378 columns



In [6]:

```
data_train.shape
```

Out[6]:

(4209, 378)

In [7]:

```
##check missing value  
data_train.isna().sum()
```

Out[7]:

```
ID      0  
y        0  
X0       0  
X1       0  
X2       0  
..  
X380     0  
X382     0  
X383     0  
X384     0  
X385     0  
Length: 378, dtype: int64
```

In [8]:

```
data_train.nunique()
```

Out[8]:

```
ID      4209  
y      2545  
X0       47  
X1       27  
X2       44  
...  
X380      2  
X382      2  
X383      2  
X384      2  
X385      2  
Length: 378, dtype: int64
```


In [9]:

```
data_train.describe()
```

Out[9]:

	ID	y	X10	X11	X12	X13	X14
count	4209.000000	4209.000000	4209.000000	4209.0	4209.000000	4209.000000	4209.000000
mean	4205.960798	100.669318	0.013305	0.0	0.075077	0.057971	0.428130
std	2437.608688	12.679381	0.114590	0.0	0.263547	0.233716	0.494865
min	0.000000	72.110000	0.000000	0.0	0.000000	0.000000	0.000000
25%	2095.000000	90.820000	0.000000	0.0	0.000000	0.000000	0.000000
50%	4220.000000	99.150000	0.000000	0.0	0.000000	0.000000	0.000000
75%	6314.000000	109.010000	0.000000	0.0	0.000000	0.000000	1.000000
max	8417.000000	265.320000	1.000000	0.0	1.000000	1.000000	1.000000

8 rows × 370 columns



In [11]:

```
## try find variance of column
print(data_train.var)
```

```
<bound method DataFrame.var of          ID          y  X0 X1  X2 X3 X4  X5 X6
X8 ... X375 X376 X377 X378 \
0      0 130.81 k v at a d u j o ... 0 0 1 0
1      6  88.53 k t av e d y l o ... 1 0 0 0
2      7  76.26 az w n c d x j x ... 0 0 0 0
3      9  80.62 az t n f d x l e ... 0 0 0 0
4     13  78.02 az v n f d h d n ... 0 0 0 0
...    ...    ... .. .. .. .. .. .. .. ..
4204 8405 107.39 ak s as c d aa d q ... 1 0 0 0
4205 8406 108.77 j o t d d aa h h ... 0 1 0 0
4206 8412 109.22 ak v r a d aa g e ... 0 0 1 0
4207 8415  87.48 al r e f d aa l u ... 0 0 0 0
4208 8417 110.85 z r ae c d aa g w ... 1 0 0 0

      X379 X380 X382 X383 X384 X385
0      0    0    0    0    0    0
1      0    0    0    0    0    0
2      0    0    1    0    0    0
3      0    0    0    0    0    0
4      0    0    0    0    0    0
...    ...    ...    ...    ...    ...
4204    0    0    0    0    0    0
4205    0    0    0    0    0    0
4206    0    0    0    0    0    0
4207    0    0    0    0    0    0
4208    0    0    0    0    0    0
```

[4209 rows x 378 columns]>

In [12]:

```
cols=[c for c in data_train.columns if 'X' in c]
print('Number of features: {}'.format(len(cols)))
```

Number of features: 376

In [13]:

```
print('Feature types:')
data_train[cols].dtypes.value_counts()
```

Feature types:

Out[13]:

```
int64      368
object       8
dtype: int64
```

In [14]:

```
##Count the data in each of the columns
counts = [[], [], []]
for c in cols:
    typ = data_train[c].dtype
    uniq = len(np.unique(data_train[c]))
    if uniq == 1:
        counts[0].append(c)
    elif uniq == 2 and typ == np.int64:
        counts[1].append(c)
    else:
        counts[2].append(c)

print('Constant features: {} Binary features: {} Categorical features: {}\n'
      .format(*[len(c) for c in counts]))
print('Constant features:', counts[0])
print('Categorical features:', counts[2])
```

Constant features: 12 Binary features: 356 Categorical features: 8

Constant features: ['X11', 'X93', 'X107', 'X233', 'X235', 'X268', 'X289',
'X290', 'X293', 'X297', 'X330', 'X347']

Categorical features: ['X0', 'X1', 'X2', 'X3', 'X4', 'X5', 'X6', 'X8']

In [16]:

```
# the prediction output
y_train = data_train['y'].values
y_train
```

Out[16]:

array([130.81, 88.53, 76.26, ..., 109.22, 87.48, 110.85])

In [17]:

```
##remove columns ID and Y from the data
remove_columns = list(set(data_train.columns) - set(['ID', 'y']))
y_train = data_train['y'].values
id_test = data_test['ID'].values

x_train = data_train[remove_columns]
x_test = data_test[remove_columns]
```

In [19]:

```
#Check for null and unique values for tain and test dataset
def check_missing_values(df):
    if df.isnull().any().any():
        print("There are missing values in the dataframe")
    else:
        print("There are no missing values in the dataframe")
```

In [20]:

```
##check null in train dataest
check_missing_values(x_train)
```

There are no missing values in the dataframe

In [21]:

```
## check null in test dataset
check_missing_values(x_test)
```

There are no missing values in the dataframe

In [22]:

```
## If for any column(s), the variance is equal to zero
## Apply Label encoder
for column in remove_columns:
    val = len(np.unique(x_train[column]))
    if val == 1:
        x_train.drop(column, axis=1) # Column with only one
        # value is useless so we drop it
        x_test.drop(column, axis=1)
    if val > 2: # Column is categorical
        mapper = lambda x: sum([ord(digit) for digit in x])
        x_train[column] = x_train[column].apply(mapper)
        x_test[column] = x_test[column].apply(mapper)
x_train.head()
```

/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:11: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

This is added back by InteractiveShellApp.init_path()
/usr/local/lib/python3.7/site-packages/ipykernel_launcher.py:12: SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame.
Try using .loc[row_indexer,col_indexer] = value instead

See the caveats in the documentation: https://pandas.pydata.org/pandas-docs/stable/user_guide/indexing.html#returning-a-view-versus-a-copy

```
if sys.path[0] == '':
```

Out[22]:

	X327	X297	X353	X382	X99	X292	X115	X257	X252	X89	...	X66	X203	X321	X19
0	1	0	0	0	0	0	0	0	0	0	...	0	0	0	
1	0	0	0	0	0	0	0	0	0	0	...	0	0	0	
2	0	0	0	1	0	0	0	0	1	0	...	0	0	0	
3	0	0	0	0	0	0	0	0	1	0	...	0	0	0	
4	0	0	0	0	0	0	0	0	1	0	...	0	0	0	

5 rows × 376 columns



In [23]:

```
print('Feature types:')
x_train[cols].dtypes.value_counts()
```

Feature types:

Out[23]:

```
int64    376
dtype: int64
```

In [24]:

```
## Perform dimensionality reduction (PCA)

from sklearn.decomposition import PCA
```

In [25]:

```
pca = PCA(n_components=12, random_state=420)
pca2_results_train = pca.fit_transform(x_train)
pca2_results_test = pca.transform(x_test)
```

In [26]:

```
## Training using xgboost
import xgboost as xgb
from sklearn.metrics import r2_score
from sklearn.model_selection import train_test_split
```

In [27]:

```
x_train, x_valid, y_train, y_valid = train_test_split(
    pca2_results_train,
    y_train, test_size=0.2,
    random_state=4242)
```

In [28]:

```
d_train = xgb.DMatrix(x_train, label=y_train)
d_valid = xgb.DMatrix(x_valid, label=y_valid)
d_test = xgb.DMatrix(pca2_results_test)
```

In [29]:

```
params = {}
params['objective'] = 'reg:linear'
params['eta'] = 0.02
params['max_depth'] = 4

def xgb_r2_score(preds, dtrain):
    labels = dtrain.get_label()
    return 'r2', r2_score(labels, preds)

watchlist = [(d_train, 'train'), (d_valid, 'valid')]
```

In [30]:

```
clf = xgb.train(params, d_train,  
               1000, watchlist, early_stopping_rounds=50,  
               feval=xgb_r2_score, maximize=True, verbose_eval=10)
```


[06:03:51] WARNING: /workspace/src/objective/regression_obj.cu:167: reg:linnear is now deprecated in favor of reg:squarederror.
[0] train-rmse:99.14835 valid-rmse:98.26297 train-r2:-58.35295
valid-r2:-67.63754
Multiple eval metrics have been passed: 'valid-r2' will be used for early stopping.

Will train until valid-r2 hasn't improved in 50 rounds.

[10]	train-rmse:81.27653	valid-rmse:80.36433	train-r2:-38.88428
	valid-r2:-44.91014		
[20]	train-rmse:66.71610	valid-rmse:65.77334	train-r2:-25.87403
	valid-r2:-29.75260		
[30]	train-rmse:54.86957	valid-rmse:53.88973	train-r2:-17.17752
	valid-r2:-19.64401		
[40]	train-rmse:45.24491	valid-rmse:44.21970	train-r2:-11.35979
	valid-r2:-12.89996		
[50]	train-rmse:37.44729	valid-rmse:36.37237	train-r2:-7.46666
	valid-r2:-8.40428		
[60]	train-rmse:31.14748	valid-rmse:30.01874	train-r2:-4.85757
	valid-r2:-5.40570		
[70]	train-rmse:26.08660	valid-rmse:24.90890	train-r2:-3.10872
	valid-r2:-3.41053		
[80]	train-rmse:22.04638	valid-rmse:20.83274	train-r2:-1.93458
	valid-r2:-2.08514		
[90]	train-rmse:18.84403	valid-rmse:17.60316	train-r2:-1.14397
	valid-r2:-1.20274		
[100]	train-rmse:16.33631	valid-rmse:15.08444	train-r2:-0.61131
	valid-r2:-0.61749		
[110]	train-rmse:14.40372	valid-rmse:13.14818	train-r2:-0.25262
	valid-r2:-0.22889		
[120]	train-rmse:12.92871	valid-rmse:11.68941	train-r2:-0.00921
	valid-r2:0.02867		
[130]	train-rmse:11.80812	valid-rmse:10.61535	train-r2:0.15815
	valid-r2:0.19897		
[140]	train-rmse:10.98603	valid-rmse:9.84998	train-r2:0.27129
	valid-r2:0.31031		
[150]	train-rmse:10.37399	valid-rmse:9.32204	train-r2:0.35023
	valid-r2:0.38226		
[160]	train-rmse:9.92029	valid-rmse:8.95919	train-r2:0.40582
	valid-r2:0.42942		
[170]	train-rmse:9.59071	valid-rmse:8.71397	train-r2:0.44464
	valid-r2:0.46022		
[180]	train-rmse:9.34334	valid-rmse:8.55560	train-r2:0.47292
	valid-r2:0.47967		
[190]	train-rmse:9.15814	valid-rmse:8.45152	train-r2:0.49361
	valid-r2:0.49225		
[200]	train-rmse:9.01373	valid-rmse:8.38985	train-r2:0.50945
	valid-r2:0.49963		
[210]	train-rmse:8.90228	valid-rmse:8.34352	train-r2:0.52151
	valid-r2:0.50514		
[220]	train-rmse:8.82529	valid-rmse:8.32079	train-r2:0.52975
	valid-r2:0.50783		
[230]	train-rmse:8.76744	valid-rmse:8.30674	train-r2:0.53589
	valid-r2:0.50950		
[240]	train-rmse:8.71781	valid-rmse:8.29981	train-r2:0.54113
	valid-r2:0.51031		
[250]	train-rmse:8.67893	valid-rmse:8.29033	train-r2:0.54522
	valid-r2:0.51143		
[260]	train-rmse:8.64604	valid-rmse:8.28552	train-r2:0.54866
	valid-r2:0.51200		
[270]	train-rmse:8.61700	valid-rmse:8.28521	train-r2:0.55168

```
valid-r2:0.51204
[280]  train-rmse:8.58766      valid-rmse:8.28510      train-r2:0.55473
valid-r2:0.51205
[290]  train-rmse:8.55980      valid-rmse:8.28708      train-r2:0.55762
valid-r2:0.51181
[300]  train-rmse:8.53510      valid-rmse:8.28718      train-r2:0.56017
valid-r2:0.51180
[310]  train-rmse:8.50959      valid-rmse:8.28848      train-r2:0.56279
valid-r2:0.51165
[320]  train-rmse:8.48800      valid-rmse:8.28924      train-r2:0.56501
valid-r2:0.51156
Stopping. Best iteration:
[272]  train-rmse:8.60985      valid-rmse:8.28187      train-r2:0.55243
valid-r2:0.51243
```

In [31]:

```
## Predict data_test values using xgboost
p_test = clf.predict(d_test)

sub = pd.DataFrame()
sub['ID'] = id_test
sub['y'] = p_test
sub.to_csv('xgb.csv', index=False)
```

In [32]:

```
p_test
```

Out[32]:

```
array([ 82.87171 , 97.357605, 83.43425 , ..., 98.78521 , 107.399536,
        96.946365], dtype=float32)
```